

The Straightness of the Calculus

Differential Topology after Milnor, and the Uniqueness of the Infinitesimal

An essay for mathematical physicists and analysts

On the plan. This essay follows the architecture of John Milnor's *Topology from the Differentiable Viewpoint* (Princeton, 1965): differentiable maps and the inverse function theorem (§1–§2), immersions and embeddings (§3), submanifolds and regular values (§4), vector fields and flows (§5), the tangent bundle (§6), and transversality and jets (§7). To that skeleton I add three chapters that the 1965 book leaves implicit but that the theme demands: connections, curvature, and the uniqueness of the Levi-Civita connection (§8); the non-uniqueness of the smooth structure — the exotic $\mathbb{R}^4$ and the exotic spheres (§9); and the failure, then repair, of the straightening in infinite dimensions — the Banach and Fréchet inverse function theorems and the Nash–Moser method (§10). A single idea runs through the whole: *the calculus is what remains when the manifold is straightened, and the question of which calculi exist, and which of them are unique, is precisely the question of how much of the straightening is canonical.*

Abstract

The differential calculus on a smooth manifold is a local linearization: a chart “straightens” the manifold to $\mathbb{R}^n$, and the tangent space, the exterior derivative, and the Lie derivative are the invariants of that straightening. This essay argues that the notion of *straightness* — the existence of a local affine model — is the organizing principle of differential topology, and that its consequences amount to a taxonomy of the possible calculi together with a precise accounting of which of them are unique (canonical) and which are not.

The first-order calculus is unique: the tangent space is the same object whether defined as equivalence classes of curves or as derivations, and the Fréchet derivative is the unique best linear approximation to a map at a point. The higher-order calculus is not unique: the k -jet for $k \geq 2$ is not determined by the 1-jet, and the spaces of connections and of metrics are genuine moduli spaces. The global calculus is obstructed: a parallelization of the tangent bundle exists only in dimensions 0, 1, 3, 7 (Adams), and the smooth structure on $\mathbb{R}^n$ is unique for $n \neq 4$ but is *not* unique in dimension 4, where there are uncountably many exotic $\mathbb{R}^4$'s. In infinite dimensions the straightening succeeds in Banach spaces (the inverse function theorem) but fails in Fréchet spaces, and is repaired by the Nash–Moser implicit function theorem — the method behind the C^∞ Nash isometric embedding theorem and the KAM theory. We close with a table of calculi and their uniqueness, and with the observation that straightness is, in the end, the notion of a local linear model, and that the “kind” of calculus a space admits is read off from how rigidly that model is determined.

0. Introduction: what “straight” means, and why it is load-bearing

A physicist who computes in flat spacetime is, without meaning to, doing differential topology. The statement “the electron follows a straight line between interactions” is an assertion about the *affine* structure of Minkowski space: there is a family of preferred curves (the geodesics of the flat connection) such that the first-order Taylor expansion of any curve about a point agrees with a linear map. The whole machinery of perturbation theory — expand about a background, keep the linear term, treat the rest as a correction — is the assertion that **locally, to first order, the world is straight.**

Differential topology is the study of what survives, and what is lost, when one tries to make that assertion on a space that is not globally straight. A smooth n -manifold M is, by

definition, a space that is *locally* $\mathbb{R}^n$ in a controlled sense: there is an atlas of charts

$$\varphi_\alpha : U_\alpha \hookrightarrow \mathbb{R}^n, \quad U_\alpha \subset M,$$

such that the transition maps $\varphi_\beta \circ \varphi_\alpha^{-1}$ are diffeomorphisms of open subsets of $\mathbb{R}^n$. A chart is a *straightening*: it replaces the curved U_α by a flat patch of $\mathbb{R}^n$. The word "diffeomorphism" does all the work — it says that the way two different straightenings agree is itself smooth, so that "straightness" is not an artifact of a particular chart but a property of the *system* of charts.

From this single sentence, three questions follow, and they are the questions of this essay:

1. **Existence.** Does a given space admit a straightening at all? (Smooth structures, embeddings, the isometric embedding problem.)
2. **Uniqueness.** If it admits one, is it unique — canonical, independent of choices? (The tangent space, the derivative, the flow, the Levi-Civita connection, the smooth structure of $\mathbb{R}^n$.)
3. **Obstruction.** If it does not admit one, or admits many, what measures the failure, or the freedom? (Characteristic classes, the hairy ball, exotic structures, moduli spaces.)

The surprising and, I want to argue, central fact is that the answer to (2) is *graded by order*. The **first-order** straightening is rigid: there is exactly one, and it is canonical. The **higher-order** straightening is flexible: there are many, and they form moduli spaces. And the **global** straightening is obstructed: it exists only in special dimensions, and in dimension four it exists in an uncountable family of mutually incompatible forms. The "kind" of calculus available on a space, and its uniqueness, is read off from this grading.

I will return to each of these. But first, the local theory.

1. The local straightening: differentiable maps and the tangent space

(Milnor, Ch. 1, §§1.1–1.2.)

Let M be a smooth n -manifold and $p \in M$. There are two standard definitions of the tangent space $T_p M$, and the fact that they agree is the first and most important uniqueness theorem of the subject.

Definition by curves. A tangent vector at p is the equivalence class of a smooth curve $\gamma : (-\varepsilon, \varepsilon) \rightarrow M$ with $\gamma(0) = p$, where $\gamma \sim \gamma'$ if, for (equivalently, for every) chart φ about p ,

$$\left. \frac{d}{dt} \right|_{t=0} \varphi(\gamma(t)) = \left. \frac{d}{dt} \right|_{t=0} \varphi(\gamma'(t)).$$

The equivalence is well-defined precisely because the transition maps are diffeomorphisms: if the two curves agree to first order in one chart, the chain rule forces them to agree in every chart. The straightening has defined an equivalence relation on curves, and the relation is independent of the straightening.

Definition by derivations. A tangent vector at p is a linear map $D : C^\infty(M) \rightarrow \mathbb{R}$ satisfying the Leibniz rule

$$D(fg) = D(f)g(p) + f(p)D(g).$$

The space of such derivations is a real vector space, and the map

$$[\gamma] \mapsto (f \mapsto \left. \frac{d}{dt} \right|_{t=0} f(\gamma(t)))$$

is a linear isomorphism $T_p M \cong \text{Der}_p(C^\infty(M))$. No choice of chart enters the second definition at all, which is why this is the definition a functional analyst reaches for: a tangent vector is a *continuous derivation* on the commutative algebra of smooth functions.

There is a third, purely algebraic, definition that makes the uniqueness transparent. Let $\mathfrak{m}_p \subset C_p^\infty(M)$ be the maximal ideal of germs of smooth functions vanishing at p . Then

$$T_p^*M \cong \mathfrak{m}_p/\mathfrak{m}_p^2,$$

the dual of the quotient of the maximal ideal by its square. The quotient $\mathfrak{m}_p/\mathfrak{m}_p^2$ is the **first-order neighborhood** of p in the algebraic sense: it is what remains of the ring of functions once one throws away everything of order ≥ 2 . That this quotient is independent of any analytic input — no metric, no connection, no chart — is the algebraic content of the statement that *the first-order structure is canonical*.

The derivative as unique best linear approximation. Let $f : M \rightarrow N$ be smooth and $p \in M$. The differential $df_p : T_pM \rightarrow T_{f(p)}N$ is characterized by the following uniqueness property: it is the **unique** linear map $L : T_pM \rightarrow T_{f(p)}N$ such that, in charts,

$$f(p+h) = f(p) + Lh + o(\|h\|) \quad (h \rightarrow 0).$$

This is the Fréchet characterization. The point for our theme is the word *unique*. Given the straightening (a pair of charts), the first-order behavior of f is captured by exactly one linear map. There is no choice, no moduli, no moduli space: the linear part of the Taylor expansion is a **functor** of the straightening, and the chain rule

$$d(g \circ f)_p = dg_{f(p)} \circ df_p$$

says that the assignment $f \mapsto df$ respects composition. The derivative is the unique natural transformation from the functor of smooth maps to the functor of linear maps between tangent spaces. In the language of the audience at hand: df is the unique natural transformation, and the chain rule is the statement of its naturality.

Two consequences are worth isolating.

- **The exterior derivative.** The 1-form df on M (for $f : M \rightarrow \mathbb{R}$) is the pullback of the standard 1-form dx on $\mathbb{R}$ under the straightening. Its coordinate-free definition, $df_p(v) = D_v f$ for $v \in T_pM$, is chart-independent for the same reason the tangent space is: it is a first-order object, and first-order objects are canonical. The full de Rham complex $\Omega^\bullet(M)$, built from df by alternation, is therefore a canonical first-order calculus. This is the calculus that a physicist uses when writing $dA = F$ for a gauge potential A : the d is the unique first-order operator that is antisymmetric in its arguments and satisfies $d^2 = 0$.
- **The Lie derivative.** For a vector field X and a tensor field T , the Lie derivative $\mathcal{L}_X T$ is the infinitesimal action of the flow of X on T . It is again a first-order object: it depends on X and T only through their values and first derivatives at a point. Its uniqueness is the statement that there is a unique derivation of the algebra of tensor fields along a given vector field that agrees with the ordinary derivative on functions. The Lie derivative is, in this sense, the unique “infinitesimal straightening of the frame” generated by X .

So at the local, first-order level, the straightening is rigid. Every object one can define — T_pM , T_p^*M , $\Omega^\bullet(M)$, $\mathcal{L}_X$, the differential — is canonical, independent of the chart, and unique. The calculus at first order is **the** calculus, not *a* calculus.

2. The inverse function theorem: local straightening, and its uniqueness up to higher order

(Milnor, Ch. 1, §1.2, and the constant rank theorem.)

The inverse function theorem is the engine that turns the *definition* of a manifold into a *theorem* about it, and it is where the “uniqueness up to higher order” theme first appears.

Theorem (inverse function theorem). Let $f : U \subset \mathbb{R}^n \rightarrow \mathbb{R}^n$ be C^1 , and suppose $Df(p)$ is invertible. Then there are neighborhoods $V \ni p$ and $W \ni f(p)$ such that $f|_V : V \xrightarrow{\sim} W$ is a C^1 diffeomorphism.

The proof (Milnor's is a contraction argument) shows something stronger than the statement: the inverse is constructed explicitly, and its first derivative is

$$D(f^{-1})_{f(p)} = (Df_p)^{-1}.$$

That is, the first-order behavior of the inverse is *exactly* the inverse of the first-order behavior of the map. The straightening and its un-straightening agree to first order, canonically. What the theorem does **not** control is the higher-order behavior: the inverse function is determined, to all orders, by the formal inverse of the Taylor series of f , but the actual (convergent) inverse is not determined by any finite jet of f . Two functions with the same 2-jet at p can have inverses that differ to second order. The straightening is unique to first order, and only *formally* unique to higher order.

This is the germ of the theme. Let me state it as a principle that will recur:

Principle (grading by order). The straightening of a manifold at a point is *canonical* in degree 1, *formal* (determined up to a choice of coordinates) in degree ≥ 2 , and *obstructed* globally. The calculi that live in degree 1 are unique; the calculi that live in degree ≥ 2 are moduli.

The **constant rank theorem** is the inverse function theorem with the invertibility hypothesis weakened to a constant-rank hypothesis, and it is the workhorse that produces submanifolds.

Theorem (constant rank). If $f : U \subset \mathbb{R}^n \rightarrow \mathbb{R}^m$ has constant rank r near p , then there are coordinates in which

$$f(x_1, \dots, x_n) = (x_1, \dots, x_r, 0, \dots, 0).$$

In words: *locally, a constant-rank map is straight*. This is the strongest local form of the theme. A map of constant rank is, after a change of coordinates on source and target (i.e., after re-straightening), exactly a projection. The "curvature" of a constant-rank map is zero, locally. What the theorem does not say — and what the rest of the essay is about — is that the re-straightening is not unique, and that the failure of global straightness is measured by invariants (curvature, characteristic classes) that are invisible to the constant rank theorem.

The inverse function theorem has a global cousin that a functional analyst will recognize: **Palais's global inverse function theorem**, which gives conditions under which a local diffeomorphism that is a proper map (equivalently, a covering map with suitable hypotheses) is a global diffeomorphism. The moral is the same: local straightening, plus a global topological hypothesis (properness, simply-connectedness of the target), upgrades to global straightening. The gap between the local and the global is exactly the gap that characteristic classes measure.

3. Immersions and embeddings: straightening the *map*, not the *space*

(Milnor, Ch. 2, §§2.1–2.3.)

Up to now the straightening has been of the *space*: a chart straightens a neighborhood in M . The next move — and the one that most directly answers a physicist's question, "can I put my curved thing into flat space?" — is to straighten a *map*.

Definition. A smooth map $f : M^n \rightarrow N^m$ is an **immersion** if df_x is injective for every $x \in M$; it is an **embedding** if, in addition, it is a homeomorphism onto its image $f(M) \subset N$.

An immersion says: *at every point, the tangent space of M sits inside the tangent space of N without collapsing*. An embedding says: *and moreover the global picture has no self-intersections*. The distinction is the distinction between “locally straight” and “globally straight”: an immersion is a local straightening of the map (the image is locally a submanifold of N), while an embedding is a global one (the image is a submanifold everywhere, with no self-contact).

The local content is again a straightening theorem:

Theorem (local form of an immersion). If $f : M^n \rightarrow N^m$ is an immersion, then for each $x \in M$ there are charts of M and N in which

$$f(y_1, \dots, y_n) = (y_1, \dots, y_n, 0, \dots, 0).$$

That is, *locally, an immersion is a straight linear inclusion*. The proof is the inverse function theorem applied to a coordinate projection of f . So an immersion is, locally, as straight as the word “straight” allows: it is a linear subspace inclusion, up to a re-straightening of source and target.

The global content is where the freedom reappears, and where the moduli live. The **hairy ball theorem** and its relatives show that the re-straightening cannot be done globally in general:

Theorem (hairy ball). S^2 admits no nonvanishing continuous tangent vector field.

In the language of straightening: there is no global nonvanishing “direction of straightness” on S^2 . The tangent bundle TS^2 is not trivial, so one cannot choose, continuously and everywhere, a parallelization — a global frame in which the manifold looks, in each direction, like $\mathbb{R}^2$. The Euler class $e(TS^2) \in H^2(S^2; \mathbb{Z}) \cong \mathbb{Z}$ evaluates to ± 2 (the Euler characteristic), which is the precise obstruction. The straightening exists locally (charts), but the *global frame* in which the manifold would be uniformly straight does not exist. This is the first appearance of the graded-uniqueness theme at the global level: the local straightening is canonical, the global straightening is obstructed, and the obstruction is a characteristic class.

The **embedding** theorems measure how much flat space is needed to straighten a given manifold globally.

Theorem (Whitney embedding). Every closed smooth n -manifold embeds in $\mathbb{R}^{2n}$, and immerses in $\mathbb{R}^{2n-1}$.

Theorem (general position / weak Whitney). For $n \geq 3$, a generic smooth map $M^n \rightarrow \mathbb{R}^{2n-1}$ is an immersion, and a generic map $M^n \rightarrow \mathbb{R}^{2n}$ is an embedding.

The numbers $2n$ and $2n - 1$ are not deep invariants; they are the dimension of the space of jets needed to make self-intersections “generic” and hence removable. What is deep is the question of whether a given immersion is *isotopic* to an embedding, and how many embeddings of a fixed isotopy class exist. That is a moduli question, and the moduli are governed by the **Smale-Hirsch theorem**:

Theorem (Smale-Hirsch). The map

$$\{\text{immersions } M^n \rightarrow N^m\} / \text{isotopy} \longrightarrow \{\text{monomorphisms } TM \hookrightarrow TN\} / \text{homotopy}$$

is a weak homotopy equivalence (for $m > n$).

In words: *an immersion is determined up to isotopy by its linear part — by the bundle monomorphism $df : TM \hookrightarrow f^*TN$* . This is the straightening principle at the level of maps: the higher-order (the actual image) carries no information beyond the first-order (the derivative), as long as one works up to isotopy. The calculus of immersions is, in this sense, a first-order calculus, and it is unique. The moment one asks for the actual geometric image (up to ambient isotopy, or with a fixed metric), the higher-order information reappears, and the moduli open up.

4. Submanifolds, regular values, and transversality: the straightening that is *defined* by a map

(Milnor, Ch. 3, §§3.1–3.3; transversality in Ch. 4.)

A submanifold is the most common object a physicist encounters as “straight inside a curved ambient space”: the world-sheet of a string in spacetime, the constraint surface in phase space, the critical manifold of a variational problem. The clean definition is via the **regular value theorem**.

Theorem (regular value / preimage of a regular value is a submanifold). Let $f : M^n \rightarrow \mathbb{R}^k$ be smooth and $q \in \mathbb{R}^k$ a regular value (i.e., df_x is surjective for every $x \in f^{-1}(q)$). Then $f^{-1}(q)$ is a closed embedded submanifold of M of dimension $n - k$.

The proof is the constant rank theorem: f has constant (maximal) rank on $f^{-1}(q)$, so locally f is a projection, and the preimage of a point is a linear subspace. The submanifold is, locally, *exactly* a straight subspace of $\mathbb{R}^n$. This is the sense in which “submanifold” and “locally straight” are the same notion: a submanifold is a subset that is, in a neighborhood of each of its points, the zero set of a submersion — i.e., a linear subspace up to straightening.

The **Sard theorem** is the analytic heart that makes regular values abundant:

Theorem (Sard). The set of critical values of a smooth $f : M^n \rightarrow \mathbb{R}^k$ is a Lebesgue null set (in fact, has measure zero in a strong sense).

So for a generic target value, the preimage is a submanifold. In the straightening language: *a generic hyperplane section of a manifold is a straight submanifold*. This is the theorem that justifies, in physics, the practice of “restricting to a generic slice” — for instance, taking a generic time-slice of a field configuration and asserting it is a smooth submanifold of configuration space.

Transversality is the relative version, and it is the workhorse of global differential topology.

Definition. A map $f : M \rightarrow N$ is **transverse** to a submanifold $S \subset N$, written $f \pitchfork S$, if for every x with $f(x) \in S$,

$$df_x(T_x M) + T_{f(x)} S = T_{f(x)} N.$$

In words: the image of the tangent space of M and the tangent space of S together span the ambient tangent space. The map f “meets” S as straightly as possible — not tangentially, but at the maximal angle allowed.

Theorem (transversality / Thom). If $f \pitchfork S$, then $f^{-1}(S)$ is a submanifold of M of codimension equal to that of S in N . Moreover, the class of maps transverse to a fixed S is dense in $C^\infty(M, N)$ (in the Whitney C^∞ topology).

The density statement is the global straightening result: *a generic map is as straight as possible relative to any prescribed submanifold*. The consequence that makes transversality the engine of the subject is the **Thom transversality theorem**, which says that a generic *family* of maps is transverse to a fixed submanifold of the jet space. This is the theorem behind the existence of generic Morse functions, generic immersions, generic vector fields, and — in physics — behind the assertion that a generic perturbation of a potential has only non-degenerate critical points, so that the critical point set is a discrete, countable, and stably defined object.

The **Morse lemma** is the local straightening of a critical point, and it is the sharpest form of the first-order rigidity:

Theorem (Morse lemma). If $f : M^n \rightarrow \mathbb{R}$ has a non-degenerate critical point at p (i.e., the Hessian $d^2 f_p$ is nondegenerate), then in some chart centered at p ,

$$f = -y_1^2 - \cdots - y_k^2 + y_{k+1}^2 + \cdots + y_n^2.$$

That is, *locally, a non-degenerate critical point is a straight quadratic form, and the only invariant is the index k* . The straightening is canonical up to an orthogonal change of coordinates — i.e., up to the action of $O(n)$ on the quadratic form. The higher-order terms of f are *completely invisible* to the local model: the Morse lemma erases them. This is the purest example of the grading: the second-order behavior is canonical (the index), and the third-and-higher-order behavior is irrelevant to the local type.

For a physicist: the Morse lemma is the mathematical content of the *stationary phase* and *saddle point* approximations. The leading term in the asymptotic expansion of

$$\int e^{i\lambda f(x)} a(x) dx$$

as $\lambda \rightarrow \infty$ is determined by the Hessian at the critical points — the straight quadratic — and the index k determines the phase factor $e^{i\pi k/2}$. The higher-order terms contribute only to the subleading coefficients. The calculus of asymptotics is, at its leading order, a first-order (well, second-order, in f) calculus, and it is unique.

5. Vector fields, flows, and the uniqueness of the infinitesimal straightening

(Milnor, Ch. 3, §3.4, and the flow box theorem.)

A vector field X on M is a choice, for each $p \in M$, of a tangent vector $X_p \in T_p M$, varying smoothly. A vector field is, pointwise, a “direction of straightness”: it tells you which direction, at each point, is the preferred direction of motion. The **flow** of X is the family of diffeomorphisms $\{\varphi_t\}_{t \in \mathbb{R}}$ generated by integrating X , satisfying $\frac{d}{dt}\varphi_t(p) = X_{\varphi_t(p)}$ and $\varphi_0 = \text{id}$.

The **flow box theorem** is the straightening result for vector fields:

Theorem (flow box). Let X be a nonvanishing vector field and $p \in M$. Then there is a chart $\varphi : U \rightarrow \mathbb{R}^n$ about p in which X is the constant vector field $\partial/\partial x^1$.

In words: *locally, a nonvanishing vector field is a straight line field, and its flow is straight translation in the x^1 direction*. The straightening is unique up to the action of the group of diffeomorphisms that fix the x^1 -direction — i.e., up to the “vertical” re-straightening of the transverse variables. So the flow is locally as straight as possible, and the only freedom in the straightening is the freedom to reparametrize the transverse coordinates.

The **uniqueness of the flow** is the key fact: given X , the flow $\{\varphi_t\}$ is the *unique* one-parameter group of diffeomorphisms with $\varphi'_t(0) = X$. This is the uniqueness theorem for the first-order linear ODE $\dot{\gamma} = X_\gamma$ with the group constraint. It is the dynamical content of the straightening: the vector field determines the flow, and the flow determines the vector field (as its infinitesimal generator), and this correspondence is bijective. There is no moduli: the flow is the unique “integration” of the straightening.

For a physicist, this is the content of the statement that a Hamiltonian vector field determines a Hamiltonian flow, and that the flow is the unique symplectomorphism group with that generator. The **Liouville equation** and the preservation of the symplectic form (and hence of volume, by the Darboux theorem) are consequences of the straightening being canonical.

The **Darboux theorem** is the symplectic straightening, and it is the sharpest possible uniqueness:

Theorem (Darboux). Every $2n$ -dimensional symplectic manifold (M, ω) is, locally, diffeomorphic to $(\mathbb{R}^{2n}, \omega_0 = \sum_{i=1}^n dx^i \wedge dx^{n+i})$.

In words: *locally, every symplectic form is the standard flat form, and the straightening is canonical*. This is the symplectic analogue of the fact that every Riemannian metric is, at

a point, the Euclidean metric — but stronger: for the symplectic form, the straightening holds on an open set, not just at a point. The reason is that the symplectic form has no local invariants (no curvature, in the Riemannian sense); its first jet determines it completely. This is the purest form of the first-order rigidity: the symplectic structure is a first-order object, and it is locally unique.

The contrast with the Riemannian case, which I will develop in §8, is the central contrast of the essay: the symplectic straightening is locally unique (no local invariants), while the Riemannian straightening is locally unique only at a point (the curvature is the first local invariant, and it appears at second order).

6. The tangent bundle: the straightening as a *bundle*

(Milnor, Ch. 4, §§4.1–4.3.)

So far the straightening has been pointwise: $T_p M$ at each p . The **tangent bundle**

$$TM = \bigsqcup_{p \in M} T_p M, \quad \pi : TM \rightarrow M$$

is the *total space* of the straightening: it is the manifold that records, over every point of M , the whole linear space in which M is straight at that point. TM is itself a smooth $2n$ -manifold, and the projection π is a submersion whose fiber over p is the linear space $T_p M \cong \mathbb{R}^n$.

The straightening theme now takes a new form: **is TM trivial?** That is, is there a global frame — n everywhere-independent vector fields $X_1, \dots, X_n$ such that $\{X_1(p), \dots, X_n(p)\}$ is a basis of $T_p M$ for every p ? If so, $TM \cong M \times \mathbb{R}^n$, and the manifold is **parallelizable**: it is, globally, as straight as a product. The answer is the first major global straightening theorem, and it is a theorem about *dimensions*:

Theorem (Adams, 1962). S^n is parallelizable if and only if $n \in \{0, 1, 3, 7\}$.

The proof uses the K -theory of the Hopf fibration and the fact that a parallelization of S^n is equivalent to the existence of n linearly independent vector fields on S^n ; the answer is read off from the Hopf invariant, which is non-zero only in dimensions 1, 3, 7 (the dimensions of the normed division algebras $\mathbb{R}, \mathbb{C}, \mathbb{H}, \mathbb{O}$). The straightening of the sphere is global if and only if the sphere is one of the three Hopf spheres. This is the sharpest possible statement of the theme at the global level: *the kind of global straightening a space admits is a function of its dimension, and the function is as rigid as the existence of division algebras.*

The **Euler class** and the **Pontryagin classes** are the obstructions to parallelization in general. For an oriented rank- n bundle $E \rightarrow M$, the Euler class $e(E) \in H^n(M; \mathbb{Z})$ is the obstruction to the existence of a nonvanishing section; the Pontryagin classes $p_i(E) \in H^{4i}(M; \mathbb{Z})$ are the obstructions to the existence of a complex structure on the real bundle. These are the invariants that measure, in cohomology, the failure of the global straightening. For a physicist: the Euler class of TS^2 is the index of the Riemann–Roch formula for the Hopf line bundle, and the Pontryagin classes are the indices that appear in the anomaly polynomial of gauge theory.

The **tangent bundle of a product** is the model case: if $M = M_1 \times M_2$, then $TM = TM_1 \times TM_2$ (pulled back to the product), and the straightening of the product is the product of the straightenings. This is why the calculus on a product is the tensor product of the calculi on the factors: the de Rham complex satisfies $\Omega^\bullet(M_1 \times M_2) \cong \Omega^\bullet(M_1) \otimes \Omega^\bullet(M_2)$, and the Euler characteristic is multiplicative, $\chi(M_1 \times M_2) = \chi(M_1)\chi(M_2)$.

7. Transversality and jets: the straightening, and its *failure to be unique*, measured by the jet

(Milnor, Ch. 4, §4.4; the jet bundle, §7 of the present essay.)

The **k -jet** of a smooth function $f : M \rightarrow \mathbb{R}$ at a point p is the equivalence class of f under the relation “agree up to order k at p .” Formally,

$$j_p^k f = f / \mathfrak{m}_p^{k+1},$$

the image of f in the quotient of the ring of germs by the $(k+1)$ -st power of the maximal ideal. The 1-jet is the tangent space (up to dualization); the 2-jet is the Hessian (plus the gradient); and so on.

The **jet bundle** $J^k(M, \mathbb{R})$ is the bundle whose fiber over p is the space of k -jets at p . It is a finite-dimensional vector bundle of rank $\binom{n+k}{k}$, and the **jet prolongation** $j^k : M \rightarrow J^k(M, \mathbb{R})$ is a smooth section. The **Thom transversality theorem**, in its full form, says that a generic function has a k -jet section that is transverse to the subbundle of degenerate jets — i.e., a generic function has non-degenerate critical points (for $k = 2$), and more generally a generic function has a k -jet that avoids a prescribed submanifold of the jet bundle.

The **non-uniqueness of the higher-order straightening** is now precise. The 1-jet of f at p is a single point of T_p^*M ; it is unique, canonical, and determined by the germ of f . The 2-jet of f at p is a point of the symmetric square $S^2 T_p^*M \oplus T_p^*M$; it is *not* determined by the 1-jet, and two functions with the same 1-jet can have different 2-jets. The k -jet for $k \geq 2$ is a *genuine invariant* of the germ, and the space of k -jets is a *moduli space* of higher-order straightenings.

This is the precise sense in which the calculus is “graded”: the 1-jet is a point (unique), the 2-jet is a vector (a choice of quadratic form), and the k -jet for $k \geq 3$ is a tensor of rank k (a choice of k -linear form). The higher the order, the more freedom, and the more the moduli space grows. The **formal inverse** of a function (the inverse of its Taylor series) is a formal power series in the k -jets, and it is unique as a formal object — but it is not determined by any finite jet, and it does not always converge. The higher-order straightening is *formally* unique (the formal inverse is unique) but *not analytically* unique (the convergent inverse is not determined by any finite jet).

The **Morse-Bott theorem** is the higher-order generalization of the Morse lemma, and it makes the grading explicit:

Theorem (Morse-Bott). If $f : M \rightarrow \mathbb{R}$ has a critical submanifold C (i.e., the set of critical points is a submanifold, and the Hessian is nondegenerate in the normal directions), then in a tubular neighborhood of C ,

$$f = f|_C + \sum_{i=1}^k y_i^2 - \sum_{j=1}^{k'} z_j^2,$$

where the y_i are normal coordinates and the z_j are also normal coordinates. The only invariant is the index k' of the normal Hessian. The higher-order terms are again invisible: the Morse-Bott lemma erases them, and the local model is a straight quadratic form on the normal bundle.

The **Morse index** and the **Morse homology** are the global invariants built from the higher-order straightening. The Morse homology $MH_*(M)$ is the homology of the chain complex generated by the critical points of a generic Morse function f , with differential counting the gradient flow lines between critical points of adjacent index. The **Morse-Weyl theorem** says that $MH_*(M) \cong H_*(M; \mathbb{Z})$ (for a generic f). This is the global straightening result: the higher-order calculus (the Morse complex) computes the same thing as the

global topology (the singular homology), and the straightening is canonical up to chain homotopy.

For a physicist, this is the mathematical content of the *instanton* and *bounce* calculations in quantum field theory: the leading contribution to the path integral comes from the critical points of the action (the classical solutions), and the index of the Hessian (the number of negative modes) determines the fermion determinant and the θ -angle dependence. The higher-order terms (the fluctuations around the classical solution) contribute to the prefactor, but the leading order is determined by the straight quadratic model.

8. Connections, curvature, and the uniqueness of the Levi-Civita connection: the second-order straightening

(Beyond Milnor; the chapter the 1965 book leaves implicit.)

Milnor's book stops at the level of the tangent bundle and the jet bundle. But the straightening theme demands one more chapter, because it is the chapter that answers the physicist's most direct question: *given a metric, is the connection unique?* The answer is **yes, at the level of the connection, and no, at the level of the metric**, and the distinction is the sharpest possible statement of the grading.

Definition. A **connection** on a vector bundle $E \rightarrow M$ is a $\mathbb{R}$ -linear map $\nabla : \Gamma(E) \rightarrow \Gamma(T^*M \otimes E)$ satisfying the Leibniz rule $\nabla(fs) = df \otimes s + f\nabla s$. In a local frame, ∇ is given by a 1-form with values in the Lie algebra of the structure group: $\nabla = d + A$, where $A \in \Omega^1(M; \mathfrak{g})$ is the **connection form**.

The **space of connections** on a fixed bundle E is an affine space modeled on $\Omega^1(M; \text{ad } E)$. That is, there is no canonical connection: the connections are in bijection with $\Omega^1(M; \text{ad } E)$, which is an infinite-dimensional vector space. The straightening at the level of the connection is *not unique*: there is a moduli space of connections, and the moduli space is the affine space of $\mathfrak{g}$ -valued 1-forms.

But now impose a metric. Let (M, g) be a Riemannian (or pseudo-Riemannian) manifold. A connection ∇ is **metric-compatible** if $\nabla g = 0$ (i.e., the connection preserves the metric: $\nabla_X g(Y, Z) = g(\nabla_X Y, Z) + g(Y, \nabla_X Z)$). A connection is **torsion-free** if $\nabla_X Y - \nabla_Y X = [X, Y]$ (i.e., the connection is symmetric in its arguments).

Theorem (Levi-Civita). There is a **unique** connection on TM that is both metric-compatible and torsion-free. It is given, in local coordinates, by the **Christoffel symbols**

$$\Gamma_{ij}^k = \frac{1}{2} g^{k\ell} (\partial_i g_{j\ell} + \partial_j g_{i\ell} - \partial_\ell g_{ij}).$$

The proof is a direct computation: the two conditions (metric-compatibility and torsion-freeness) give two linear equations for the three unknowns Γ_{ij}^k , Γ_{ji}^k , and the antisymmetric part of Γ_{ij}^k , and the unique solution is the Christoffel formula. The straightening is unique: *given the metric, there is exactly one connection that straightens the manifold in the metric sense (parallel transport preserves lengths and angles) and the symmetry sense (no torsion).*

This is the central uniqueness theorem of the essay, and it is the answer to the physicist's question. The **Riemannian metric** is a choice (a section of $S^2 T^*M$), and the space of metrics is an infinite-dimensional convex cone. But **given a metric, the connection is unique**: the Levi-Civita connection is the unique metric-compatible, torsion-free connection. The straightening at the level of the connection is canonical, once the metric is fixed. The freedom has moved up one level: it is not in the connection (which is unique) but in the metric (which is a moduli space).

The **curvature** $R^\nabla \in \Omega^2(M; \text{End}(TM))$ is the obstruction to the connection being flat, and it is the first non-vanishing local invariant of the metric, appearing at second order. In

a normal coordinate system at a point p (geodesic coordinates), $g_{ij}(p) = \delta_{ij}$ and $\Gamma_{ij}^k(p) = 0$, but $\partial_k g_{ij}(p) = 0$ while $\partial_k \partial_\ell g_{ij}(p) \neq 0$ in general, and the curvature is

$$R_{jkl}^i(p) = \partial_k \Gamma_{jl}^i(p) - \partial_l \Gamma_{jk}^i(p) + \Gamma_{mk}^i(p) \Gamma_{jl}^m(p) - \Gamma_{ml}^i(p) \Gamma_{jk}^m(p) = \partial_k \partial_l g_{jk}(p) - \partial_k \partial_j g_{lk}(p) + \dots$$

The curvature is a second-order invariant: it is invisible to the first jet of g (which is zero at p in normal coordinates), but visible to the second jet. This is the Riemannian analogue of the symplectic contrast: the symplectic form has no local invariants (its first jet determines it), while the Riemannian metric has local invariants starting at second order (the curvature). The straightening of the Riemannian manifold is unique at a point (the metric is Euclidean, the connection is flat, in normal coordinates), but not on an open set (the curvature is the obstruction).

The **Ricci tensor** and the **scalar curvature** are the contractions of the curvature, and they are the invariants that appear in the Einstein-Hilbert action

$$S_{\text{EH}}[g] = \int_M R \text{vol}_g.$$

The **Einstein equations** $G_{ij} + \Lambda g_{ij} = 8\pi T_{ij}$ are the critical point equations of this action, and they are, in the language of the straightening, the equations that the metric must satisfy for the straightening to be "as flat as possible" in the presence of matter. The uniqueness of the Levi-Civita connection is the mathematical content of the statement that *given the metric, the geodesics are unique*: the straight lines of the curved space are the geodesics of the unique connection, and they are the unique curves that are "as straight as possible" in the metric sense.

The **Nash embedding theorem** is the global straightening result for Riemannian manifolds:

Theorem (Nash, 1956, C^k ; 1965, C^∞ via Nash-Moser). Every n -dimensional Riemannian manifold (M, g) is isometrically embeddable in $\mathbb{R}^N$ for N sufficiently large. Nash's original 1956 result gives a C^k isometric embedding ($k \geq 1$) for $N = n(3n + 11)/2$, and his 1965 C^∞ result (which requires the Nash-Moser method, §10) achieves the same dimension $N = n(3n + 11)/2$ in the C^∞ category. (Improved, and in some regimes optimal, bounds are known, but the precise sharp value of N for C^∞ isometric embedding is a separate and largely open question.)

In words: *every curved space can be straightened, globally, by embedding it in a sufficiently high-dimensional flat space*. The isometric embedding is the global straightening: the metric g is the restriction of the Euclidean metric on $\mathbb{R}^N$ to the image of M . The straightening is unique up to rigid motions of $\mathbb{R}^N$ (the isometry group), but the embedding itself is not unique: there are many isometric embeddings of the same (M, g) into $\mathbb{R}^N$, and the moduli space of embeddings is a genuine moduli space. The C^∞ case requires the Nash-Moser method, which I will describe in §10.

The **Bianchi identities** $\nabla_{[a} R_{bc]de} = 0$ and $\nabla^a R_{abcd} = \nabla_c R_{bd} - \nabla_d R_{bc}$ are the integrability conditions for the curvature, and they are the statement that the curvature is "the derivative of the connection" in a sense that is compatible with the Leibniz rule. They are the higher-order integrability conditions that the straightening must satisfy, and they are the reason the curvature is a tensor (coordinate-independent) rather than just a collection of coordinate-dependent quantities.

9. The non-uniqueness of the smooth structure: exotic $\mathbb{R}^4$ and the exotic spheres

(Beyond Milnor; the chapter the 1965 book cannot contain.)

Milnor's book is written in the C^∞ category, and it assumes that the smooth structure on $\mathbb{R}^n$ is unique. This is true for $n \neq 4$, and it is **false** for $n = 4$. The failure of uniqueness in dimension four is the most dramatic and most mysterious of all the non-uniqueness results in differential topology, and it is the result that most directly answers the question "is the straightening unique?" with a resounding **no, in dimension four**.

Definition. A **smooth structure** on a topological n -manifold M is an equivalence class of C^∞ atlases, where two atlases are equivalent if the union is a smooth atlas. Two smooth structures on the same topological manifold are **exotic** if they are not diffeomorphic (even though they are homeomorphic).

Theorem (Mazur, 1966; Freedman, 1982; Donaldson, 1983). There exist exotic $\mathbb{R}^4$'s: smooth structures on the topological 4-manifold $\mathbb{R}^4$ that are not diffeomorphic to the standard $\mathbb{R}^4$. Moreover, there are **uncountably many** such structures.

The construction of an exotic $\mathbb{R}^4$ proceeds as follows. Mazur constructed a compact contractible 4-manifold W (the **Mazur manifold**) with ∂W a homology 3-sphere that is not simply connected. The interior of W is diffeomorphic to $\mathbb{R}^4$ minus a point. But the boundary ∂W is not S^3 , so W cannot be the standard 4-ball B^4 . Gluing a copy of B^4 to W along ∂W (via a homeomorphism, which exists by the h-cobordism theorem in dimension 3) gives a smooth structure on $\mathbb{R}^4$ that is not the standard one. The non-diffeomorphism is detected by the **Donaldson invariants** (for definite 4-manifolds) or by the **Seiberg-Witten invariants** (for general 4-manifolds), which are sensitive to the smooth structure and not to the homeomorphism type.

The **uncountability** comes from the fact that the set of exotic $\mathbb{R}^4$'s can be constructed by "small" modifications of the standard structure in arbitrarily small regions, and the set of such modifications is uncountable (parameterized by a subset of the space of embeddings of a 4-ball into $\mathbb{R}^4$, which is uncountable). This is in stark contrast to $n \neq 4$: for $n \leq 3$, the smooth structure on $\mathbb{R}^n$ is unique (the $n = 3$ case is the Poincaré conjecture, proved by Perelman in 2003), and for $n \geq 5$, the smooth structure on $\mathbb{R}^n$ is unique by the h -cobordism theorem.

The **exotic spheres** are the higher-dimensional analogue, and they are the first example of non-uniqueness in the smooth category that Milnor himself discovered:

Theorem (Milnor, 1961). There exist smooth 7-spheres that are not diffeomorphic to the standard S^7 . In fact, the group Θ_7 of diffeomorphism classes of smooth 7-spheres (under connected sum) is isomorphic to $\mathbb{Z}/28\mathbb{Z}$.

Milnor's construction uses the **Hopf fibration** $S^3 \rightarrow S^7 \rightarrow S^4$ and the fact that the tangent bundle of S^7 is trivial (since $7 \in \{0, 1, 3, 7\}$). More generally, the exotic 7-spheres arise as the total spaces of S^3 -bundles over S^4 ; such bundles are classified by $\pi_3(SO(4)) \cong \mathbb{Z} \oplus \mathbb{Z}$, and the total spaces that are homotopy 7-spheres fall into exactly 28 diffeomorphism classes, distinguished by the **Euler number** of the bundle (equivalently, by the Milnor invariant μ , a smooth invariant invisible to the homeomorphism type). The group Θ_7 of diffeomorphism classes of smooth 7-spheres under connected sum is $\mathbb{Z}/28\mathbb{Z}$. The group Θ_n of exotic n -spheres is finite for $n \geq 5$ (its order is known for many n), and it vanishes in the sense that $\mathbb{R}^n$ has a unique smooth structure for $n \neq 4$ — but the $n = 4$ case is the exception, and the exception is uncountable.

The **4-dimensional Poincaré conjecture** (proved by Freedman in 1982, in the topological category) says that every simply-connected closed topological 4-manifold homeomorphic to S^4 is homeomorphic to S^4 . But the **smooth** 4-dimensional Poincaré conjecture — that every simply-connected closed smooth 4-manifold homeomorphic to S^4 is *diffeomorphic* to S^4 — is **open**. This is the single most important open problem in 4-dimensional differential topology, and it is a direct question about the uniqueness of the straightening: *is the smooth structure on S^4 unique?* The existence of exotic $\mathbb{R}^4$'s suggests the answer is no, but a proof is lacking.

For a physicist, the relevance is direct: the **instanton number** in Yang–Mills theory on

$\mathbb{R}^4$ is a topological invariant (the second Chern class of the associated bundle), and it is sensitive to the smooth structure. The existence of exotic $\mathbb{R}^4$'s means that the instanton moduli space depends on the smooth structure of the base, and the **Donaldson invariants** of a 4-manifold are, in physics, the partition function of the $\mathcal{N} = 2$ supersymmetric Yang-Mills theory on that 4-manifold (Witten's interpretation, 1994). The non-uniqueness of the smooth structure is, in physics, the non-uniqueness of the vacuum structure of the gauge theory on the same topological manifold.

10. The straightening in infinite dimensions: Banach, Fréchet, and Nash-Moser

(The chapter that answers the functional analyst's question.)

Everything so far has been finite-dimensional, and the straightening has been a diffeomorphism onto an open set of $\mathbb{R}^n$. The functional analyst's question is what happens when the "manifold" is an infinite-dimensional space of maps, and the "straightening" is the inverse function theorem for a nonlinear operator between such spaces. The answer is that the finite-dimensional straightening **succeeds in Banach spaces, fails in Fréchet spaces, and is repaired by the Nash-Moser method**.

The Banach inverse function theorem. Let X, Y be Banach spaces and $F : X \rightarrow Y$ a C^1 map with $F(0) = 0$. If $DF_0 : X \rightarrow Y$ is an **isomorphism** (bounded, bijective, with bounded inverse), then F is locally a diffeomorphism: there are neighborhoods $U \ni 0$ in X and $V \ni 0$ in Y such that $F|_U : U \rightarrow V$ is a C^1 diffeomorphism.

The proof is the same contraction argument as in the finite-dimensional case, and it works because the derivative DF_0 is an *isomorphism of Banach spaces* — a bounded invertible linear map. The straightening is unique to first order (the derivative is the unique best linear approximation, as in §1), and the higher-order behavior is controlled by the same formal inverse as in §2. The Banach inverse function theorem is the infinite-dimensional straightening, and it is as canonical as the finite-dimensional one: the derivative is the unique first-order object, and the local diffeomorphism is unique.

The failure in Fréchet spaces. A **Fréchet space** is a complete metrizable locally convex space that is not, in general, a Banach space (it has no norm, only a countable family of seminorms). The space $C^\infty(M)$ of smooth functions on a compact manifold M , with its standard Fréchet topology (uniform convergence of all derivatives on compact sets), is the model example. The Banach inverse function theorem **fails** for Fréchet spaces: the derivative DF_0 may be a surjective continuous linear map with a continuous right-inverse (a *splitting*), but the nonlinear map F need not be locally a diffeomorphism.

The reason is that the contraction argument requires a *norm* to control the error, and a Fréchet space has no norm — only a family of seminorms. The contraction constant, which is uniform in the finite-dimensional and Banach cases, is not uniform in the Fréchet case: the seminorms are independent, and a small error in one seminorm can be large in another. The straightening, which is uniform in the Banach case, is **not uniform** in the Fréchet case, and the local diffeomorphism can fail to exist.

The Nash-Moser implicit function theorem. The repair is the **Nash-Moser method**, due to Nash (1956, for the C^1 isometric embedding theorem) and Moser (1965, for the C^∞ isometric embedding theorem and the Kolmogorov-Arnold-Moser (KAM) theorem). The key idea is to replace the Newton iteration of the Banach inverse function theorem with a **modified Newton iteration** that includes a *smoothing operator* at each step.

The standard Newton iteration for solving $F(x) = 0$ is

$$x_{n+1} = x_n - (DF_{x_n})^{-1}F(x_n).$$

The problem in the Fréchet case is that $(DF_{x_n})^{-1}$ loses derivatives: if F is a differential operator of order k , then $(DF_{x_n})^{-1}$ gains k derivatives, and the iteration is not convergent

in the C^∞ topology (the derivative loss accumulates). The Nash-Moser fix is to insert a **smoothing operator** S_λ (a family of bounded operators, parameterized by a smoothing parameter $\lambda \rightarrow 0$, that regularizes by λ^{-k} derivatives at the cost of a factor λ^k in the norm) into the iteration:

$$x_{n+1} = x_n - S_{\lambda_n}((DF_{x_n})^{-1}F(x_n)).$$

The smoothing operator S_{λ_n} is chosen so that the derivative loss is compensated by the smoothing, and the iteration converges **quadratically** (as in the finite-dimensional case) but with a loss of a fixed number of derivatives at each step. The convergence is in the **tame Fréchet** topology (a topology slightly weaker than C^∞ , in which the smoothing operators are tame maps), and the result is a local diffeomorphism in the tame category, which is sufficient for the applications.

The C^∞ Nash isometric embedding theorem is the first and most celebrated application. The isometric embedding problem is to find a smooth map $f : M^n \rightarrow \mathbb{R}^N$ such that $f^*(\delta) = g$ (the pullback of the Euclidean metric δ is the given metric g). The relevant operator is

$$F(f) = f^*(\delta) - g \in \Gamma(S^2T^*M),$$

and the derivative $DF_f : \Gamma(TM) \rightarrow \Gamma(S^2T^*M)$ is

$$DF_f(h) = \mathcal{L}_h g = \nabla h + (\nabla h)^t - \nabla f^t h$$

(the Lie derivative of g along h , expressed in terms of the Levi-Civita connection ∇). The operator DF_f is a **first-order differential operator**, and it is surjective for N sufficiently large (the surjectivity is the linearized isometric embedding theorem, proved by Nash). But the right-inverse $(DF_f)^{-1}$ loses one derivative, and the standard Newton iteration does not converge in C^∞ . The Nash-Moser method, with its smoothing operators, gives the convergence, and the C^∞ isometric embedding theorem follows.

The KAM theorem is the second great application, and it is the one that a mathematical physicist will recognize. The KAM theorem says that a **near-integrable** Hamiltonian system — a Hamiltonian $H = H_0 + \varepsilon H_1$ with H_0 integrable and ε small — has, for ε sufficiently small, a **Cantor set** of invariant tori on which the motion is quasi-periodic. The relevant operator is the **homological equation**

$$\omega \cdot \partial_\theta u = f(\theta) - \bar{f},$$

where $\omega \in \mathbb{R}^n$ is the frequency vector, $f : \mathbb{T}^n \rightarrow \mathbb{R}$ is a smooth function, and $\bar{f}$ is its average. The operator $\omega \cdot \partial_\theta$ has a **small divisor problem**: its inverse is

$$(\omega \cdot \partial_\theta)^{-1} f = \sum_{k \neq 0} \frac{f_k}{i \omega \cdot k} e^{ik \cdot \theta},$$

and the denominators $\omega \cdot k$ can be arbitrarily small (the small divisors), so the inverse loses derivatives (the $1/(\omega \cdot k)$ factor grows like $|k|^{-\tau}$ for some $\tau > 0$). The standard Newton iteration does not converge. The Nash-Moser method, with its smoothing operators and a **Diophantine condition** on ω (the frequencies are "sufficiently irrational": $|\omega \cdot k| \geq c |k|^{-\tau}$ for $c > 0, \tau > 0$), gives the convergence, and the KAM theorem follows.

For a physicist, the KAM theorem is the mathematical content of the statement that *near-integrable systems are, to first order, integrable, and the perturbation theory (the Newton iteration) is valid up to a fixed number of derivatives, but the convergence requires the Diophantine condition and the smoothing operator*. The straightening of the near-integrable system is the integrable system H_0 (the flat connection, the straightening), and the perturbation H_1 is the curvature (the obstruction to the straightening being exact). The KAM theorem says that the straightening is unique up to a canonical transformation (the smoothing operator), and the uniqueness is conditional on the Diophantine condition.

The **Moser's stability theorem** and the **Nash-Moser theory of the C^∞ Nash embedding** are the two pillars of the infinite-dimensional straightening, and they are the reason that the C^∞ calculus on an infinite-dimensional space is not the same as the C^1 calculus: the C^∞ straightening requires the Nash-Moser method, and the C^1 straightening does not. The "kind" of calculus available on an infinite-dimensional space is, in this sense, graded by the regularity: the C^1 calculus is unique (the Banach inverse function theorem), and the C^∞ calculus is unique only after the Nash-Moser repair.

11. The table of calculi, and the straightness principle

We can now collect the results of the essay into a single table, organized by the order of the straightening and the uniqueness of the resulting calculus.

Object	Order	Unique?	Moduli / obstruction	Milnor reference
Tangent space $T_p M$	1	Yes (canonical)	None	Ch. 1, §1.1
Differential df_p	1	Yes (unique best linear approx.)	None	Ch. 1, §1.2
Exterior derivative d	1	Yes (unique $d^2 = 0$)	None	—
Lie derivative $\mathcal{L}_X$	1	Yes (unique derivation along X)	None	—
Inverse function (local)	1	Yes to 1st order; formal to higher	Formal inverse unique, convergent inverse not by finite jet	Ch. 1, §1.2
Flow of a vector field	1	Yes (unique 1-param. group)	None	Ch. 3, §3.4
Symplectic form (local)	1	Yes (Darboux: no local invariants)	None (locally)	—

$k - jet, k \geq 2$	k	No	$S^k T_p^* M$ (moduli of k -linear forms)	—
Morse index	2	Yes (canonical, up to $O(n)$)	$Index k \in \{0, \dots, n\}$	—
Immersion (up to isotopy)	1	Yes (Smale- Hirsch: deter- mined by df)	$Bundle monomorphism TM \hookrightarrow f^*TN$	Ch. 2, §2.2
Embedding (global)	1	No	Moduli of em- beddings (Whit- ney $2n$, $2n - 1$)	Ch. 2, §2.3
Parallelization of S^n	Global	No <i>** No ** (only $n \in \{0, 1, 3, 7\}$)</i>	Adams: K -theory of Hopf fibration	—
Riemannian metric	2	No	$Convex cone \Gamma(S^2 T^* M)$	—
Levi-Civita con- nection	2	Yes (given g)	None (given g)	—
Curvature R^∇	2	Yes (given ∇)	Obstruction to flatness	—
$Smooth structure on \mathbb{R}^n, n \neq 4$	Global	Yes (unique)	None	—

Smooth structure on $\mathbb{R}^4$	Global	No (uncountably many)	Exotic $\mathbb{R}^4$'s (Mazur, Freedman, Donaldson)	—
Exotic 7-spheres	Global	<i>** No **</i> ($\Theta_7 \cong \mathbb{Z}/28\mathbb{Z}$)	Milnor invariant μ	—
Banach inverse function theorem	1 (Banach)	Yes	None (isomorphism of Banach spaces)	—
Fréchet inverse function theorem	1 (Fréchet)	Fails	Derivative loss (no uniform norm)	—
Nash-Moser implicit function theorem	1 (tame Fréchet)	Yes (after repair)	Smoothing operators + Diophantine condition	—

The pattern is the **straightness principle**, stated in its final form:

The straightness principle. The calculus on a manifold is a local linearization, and its uniqueness is graded by order. The first-order calculus (the tangent space, the derivative, the flow, the symplectic form) is **unique and canonical**: there is exactly one, and it is independent of all choices. The higher-order calculus (the k -jets for $k \geq 2$, the metrics, the embeddings) is **not unique**: it is a moduli space, and the moduli are measured by characteristic classes, the curvature, or the smooth structure. The global calculus (the parallelization, the smooth structure of $\mathbb{R}^n$) is **obstructed**: it exists only in special dimensions, and in dimension four it exists in an uncountable family. In infinite dimensions, the first-order straightening succeeds in Banach spaces, fails in Fréchet spaces, and is repaired by the Nash-Moser method. The "kind" of calculus a space admits, and its uniqueness, is read off from this grading.

The straightness is, in the end, the notion of a **local linear model**, and the calculus is the **invariant** of that model. The first-order model is unique, so the first-order calculus is unique. The higher-order model is not unique, so the higher-order calculus is a moduli. The global model is obstructed, so the global calculus is obstructed. This is the content of the essay, and it is the content of differential topology.

For the mathematical physicist, the moral is simple and load-bearing: **when you expand about a background and keep the linear term, you are using the unique first-order calculus, and the result is independent of the background (gauge-invariant, coordinate-invariant, chart-invariant). When you keep the quadratic term, you are leaving the unique calculus and entering the moduli, and the result depends on the background (the curvature, the connection, the metric). When you try to make the straightening global, you are asking for a parallelization or a global frame, and the answer is a characteristic class (the Euler number, the Pontryagin class, the instanton number). The "straightness" of the calculus is what makes perturbation theory work, and the "curvature" is what makes it fail, and the failure is measured by an invariant that is, in physics, a topological charge.**

References

- J. Milnor, *Topology from the Differentiable Viewpoint*, Princeton University Press, 1965.
- J. Milnor, "On the existence of a connection with curvature zero," *Pacific J. Math.* **9** (1959), 64–72.
- J. Milnor, "Differentiable structures on S^7 ," *Proc. Amer. Math. Soc.* **12** (1961), 457–463.
- R. Palais, "A global form of the inverse and implicit function theorems," *Amer. J. Math.* **88** (1966), 400–422.
- M. Hirsch, *Differential Topology*, Springer, 1976.
- V. Guillemin and A. Pollack, *Differential Topology*, Prentice-Hall, 1974.
- J. M. Lee, *Introduction to Smooth Manifolds*, 2nd ed., Springer, 2013.
- M. Gromov, *Partial Differential Relations*, Springer, 1986.
- J. Nash, "The imbedding problem for Riemannian manifolds," *Ann. of Math.* **63** (1956), 20–63.
- J. Nash, "The imbedding problem for Riemannian manifolds. II," *Ann. of Math.* **78** (1963), 173–222.
- J. Moser, "On the iteration of certain nonlinear elliptic operators," *Comm. Pure Appl. Math.* **16** (1963), 565–573.
- J. Moser, "On the Poincaré recurrence theorem," *Ann. of Math.* **82** (1965), 197–212.
- S. Smale, "Differentiable dynamical systems," *Bull. Amer. Math. Soc.* **73** (1967), 747–817.
- M. Hirsch, "Immersion and embedding of differentiable manifolds," *Ann. of Math.* **74** (1961), 546–559.
- S. Donaldson, *Lectures on the Topology of 4-Manifolds*, 2nd ed., World Scientific, 1990.
- E. Witten, "Supersymmetry and Morse theory," *J. Differential Geom.* **17** (1982), 661–692.
- E. Witten, "Supersymmetry and the Atiyah–Singer index theorem," *J. Geom. Phys.* **9** (1992), 141–212.
- E. Witten, "On the structure of the topological phase of two-dimensional gravity," *Nucl. Phys. B* **340** (1990), 281–332.

End of essay.